

```
In [1]:  import pandas as pd
import numpy as np
```

```
In [4]:  df = pd.read_csv('Advertising.csv')
```

```
In [5]:  df
```

Out[5]:

	Unnamed: 0	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9
...
195	196	38.2	3.7	13.8	7.6
196	197	94.2	4.9	8.1	9.7
197	198	177.0	9.3	6.4	12.8
198	199	283.6	42.0	66.2	25.5
199	200	232.1	8.6	8.7	13.4

200 rows × 5 columns

```
In [ ]:  ## DATA PREPROCESSING ##
```

```
In [6]:  df.head()
```

Out[6]:

	Unnamed: 0	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9

In [7]: `df.tail()`

Out[7]:

	Unnamed: 0	TV	Radio	Newspaper	Sales
195	196	38.2	3.7	13.8	7.6
196	197	94.2	4.9	8.1	9.7
197	198	177.0	9.3	6.4	12.8
198	199	283.6	42.0	66.2	25.5
199	200	232.1	8.6	8.7	13.4

In [8]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Unnamed: 0      200 non-null   int64
1   TV              200 non-null   float64
2   Radio           200 non-null   float64
3   Newspaper       200 non-null   float64
4   Sales           200 non-null   float64
dtypes: float64(4), int64(1)
memory usage: 7.9 KB
```

In [9]: `df.shape`

Out[9]: (200, 5)

In [10]: `df.columns`

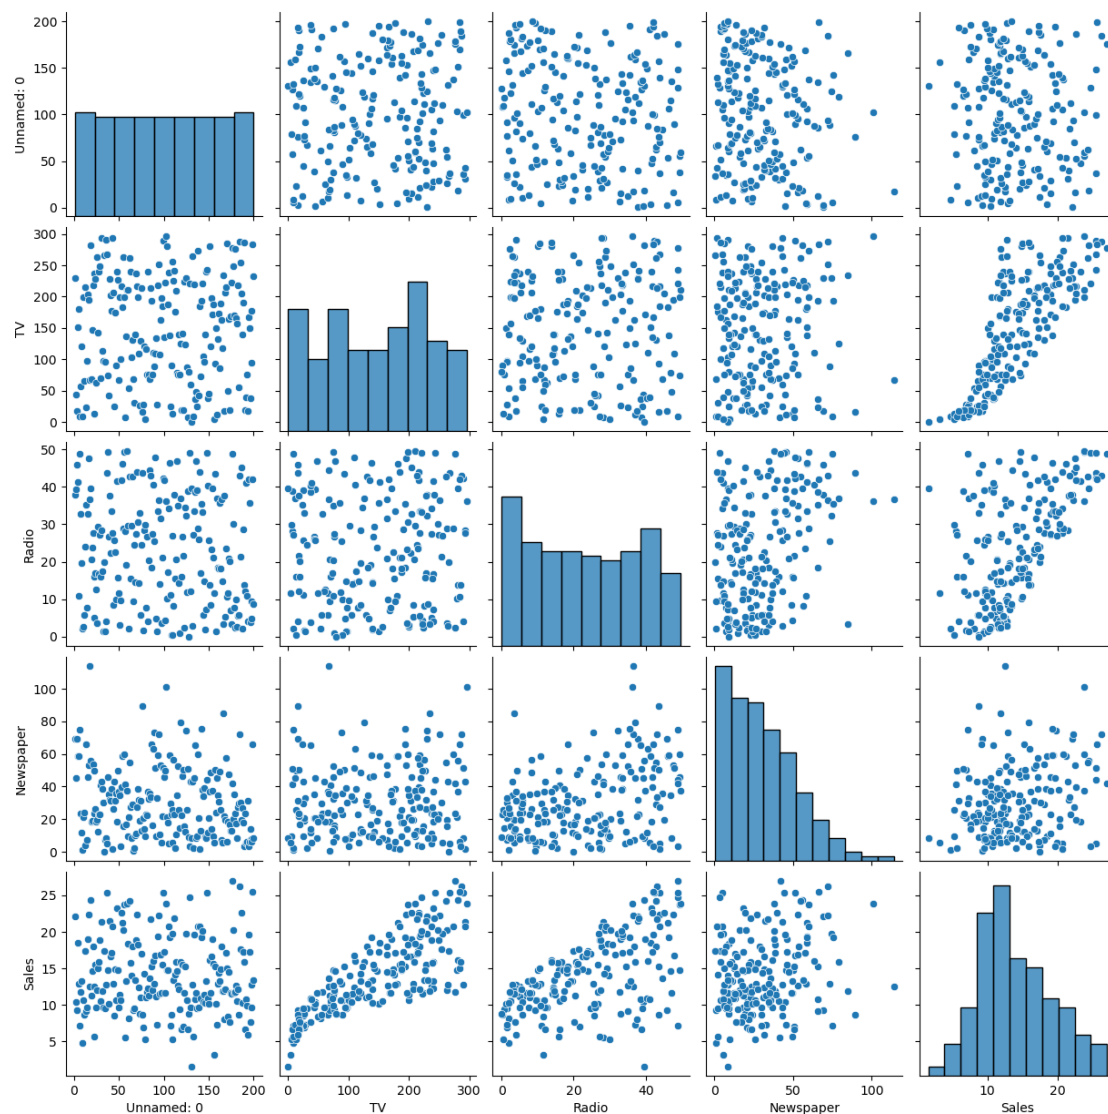
Out[10]: Index(['Unnamed: 0', 'TV', 'Radio', 'Newspaper', 'Sales'], dtype='object')

In [11]: `df.isnull().sum()`

Out[11]: Unnamed: 0 0
TV 0
Radio 0
Newspaper 0
Sales 0
dtype: int64


```
In [12]: ## VISUALIZATION ##  
  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
sns.pairplot(df)
```

Out[12]: <seaborn.axisgrid.PairGrid at 0x20ccec45b20>




In [22]:


```
sns.heatmap(df.corr(),annot=True)  
plt.show()
```




```
In [23]: ▶ sns.lmplot(data=df,x='Radio',y="Sales")  
plt.show()
```



```
In [24]: ▶ sns.lmplot(data=df,x='TV',y="Sales")  
plt.show()
```



```
In [25]: ▶ sns.lmplot(data=df,x='Newspaper',y="Sales")  
plt.show()
```



In [29]: `## Splitting the data ##`

```
x=df[['TV','Radio','Newspaper']]
y=df['Sales']
print(len(x), len(y))
print(x)
print(y)
print(x.shape)
print(y.shape)
```

200 200

	TV	Radio	Newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4
..
195	38.2	3.7	13.8
196	94.2	4.9	8.1
197	177.0	9.3	6.4
198	283.6	42.0	66.2
199	232.1	8.6	8.7

[200 rows x 3 columns]

0	22.1
1	10.4
2	9.3
3	18.5
4	12.9
...	...
195	7.6
196	9.7
197	12.8
198	25.5
199	13.4

Name: Sales, Length: 200, dtype: float64
(200, 3)
(200,)

In [30]: `## Training and Test Data ##`

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test=train_test_split(x,y,test_size=1/10,randc
```



```
In [31]:  ► ## Linear Regression ##

from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x_train,y_train)
model.score(x_test,y_test)
```

Out[31]: 0.7556846087219637

```
In [32]:  ► from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(x_train, y_train)
```

Out[32]: LinearRegression()

```
In [33]:  ► t_data_predic = regressor.predict(x_train)
```

```
In [34]:  ► ## Error Calculation ##

from sklearn import metrics
error_score = metrics.r2_score(y_train, t_data_predic)
print("R squared Error : ", error_score)
```

R squared Error : 0.9126527610435033

In [40]: `## Plotting THE data ##`

```
plt.scatter(y_train, t_data_predic)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices")
plt.show()
```



```
In [43]: ▶ ## Adding the Linear Line ##  
  
plt.scatter(y_train, t_data_predic)  
plt.xlabel("Actual Price")  
plt.ylabel("Predicted Price")  
plt.title(" Actual Prices vs Predicted Prices")  
plt.plot([min(y_train), max(y_train)], [min(y_train), max(y_train)], 'r--')  
plt.show()
```



```
In [44]: ▶ b = regressor.coef_  
print("Coefficient :",b)  
  
Coefficient : [ 0.04445735  0.19820881 -0.00245565]
```

```
In [45]: ▶ a = regressor.intercept_  
print("Intercept :",a)  
  
Intercept : 3.0013503531930663
```